



Type something...

01) SVM

01_Data Preprocessing

02_svm_loss_naive(W, x, y, reg)

03_svm_loss_vectorized(W, x, y, reg)

04_train

1. Inline Question 2

01) SVM

01_Data Preprocessing

- 데이터 각각을 표준화시켜서 전처리를 진행하였다.

```
[8] # Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('dev data shape: ', X_dev.shape)

Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

02_svm_loss_naive(W, x, y, reg)

- 함수는 총 2개를 구현하는 것이 목적이었다. 먼저 svm_loss_naive를 구현해보자.
- reg는 규제 정도로 loss를 구한 후 W의 L2 norm과 곱해서 loss에 더할 것이다.
- 먼저 loss를 구하는 것이 목적이었는데 각각의 loss를 구하기 위해서 먼저 score를 계산해주었다.
- train data와 W를 (행렬곱)시켜줘서 값을 계산해준 후 y[i]에 해당하는 correct score를 따로 선언하였다.
- 그 후 score - correct_score + 1과 0의 max 값을 할당해주었다. 이와 동시에 dW를 계산해주어야 했는데 이는 ppt의 나와있는 식을 참고하여 계산하였다.
- 평균을 구해줘야하기 때문에 모든 열에다가 다 더해준 후 num_train으로 나눠줬다.

Computing the gradient analytically with Calculus

The numerical gradient is very simple to compute using the finite difference approximation, but the downside is that it is approximate (since we have to pick a small value of h , while the true gradient is defined as the limit as h goes to zero), and that it is very computationally expensive to compute. The second way to compute the gradient is analytically using Calculus, which allows us to derive a direct formula for the gradient (no approximations) that is also very fast to compute. However, unlike the numerical gradient it can be more error prone to implement, which is why in practice it is very common to compute the analytic gradient and compare it to the numerical gradient to check the correctness of your implementation. This is called a **gradient check**.

Lets use the example of the SVM loss function for a single datapoint:

$$L_i = \sum_{j \neq y_i} \left[\max(0, w_j^T x_i - w_{y_i}^T x_i + \Delta) \right]$$

We can differentiate the function with respect to the weights. For example, taking the gradient with respect to w_{y_i} we obtain:

$$\nabla_{w_{y_i}} L_i = - \left(\sum_{j \neq y_i} \mathbf{1}(w_j^T x_i - w_{y_i}^T x_i + \Delta > 0) \right) x_i$$



where $\mathbf{1}$ is the indicator function that is one if the condition inside is true or zero otherwise. While the expression may look scary when it is written out, when you're implementing this in code you'd simply count the number of classes that didn't meet the desired margin (and hence contributed to the loss function) and then the data vector \mathbf{x}_i scaled by this number is the gradient. Notice that this is the gradient only with respect to the row of \mathbf{W} that corresponds to the correct class. For the other rows where $j \neq y_i$ the gradient is:

$$\nabla_{w_j} L_i = \mathbf{1}(w_j^T \mathbf{x}_i - w_{y_i}^T \mathbf{x}_i + \Delta > 0) \mathbf{x}_i$$

```
def svm_loss_naive(W, X, y, reg):
    """
    Structured SVM loss function, naive implementation (with loops).

    Inputs have dimension D, there are C classes, and we operate on minibatches
    of N examples.

    Inputs:
    - W: A numpy array of shape (D, C) containing weights.
    - X: A numpy array of shape (N, D) containing a minibatch of data.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c means
        | that X[i] has label c, where 0 <= c < C.
    - reg: (float) regularization strength

    Returns a tuple of:
    - loss as single float
    - gradient with respect to weights W; an array of same shape as W
    """
    dW = np.zeros(W.shape) # initialize the gradient as zero

    # compute the loss and the gradient
    num_classes = W.shape[1] # 10
    num_train = X.shape[0] # X의 개수 -> 500, 49000....
    loss = 0.0
    #print(X[0].shape) # 1x3073
    #print(W.shape) # 3073 x 10
    #print(num_train) # 500
    for i in range(num_train):
        scores = X[i].dot(W) # X가 전치행렬로 바꿔어서 계산
        correct_class_score = scores[y[i]] # np.dot와 같음
        for j in range(num_classes):
            if j == y[i]:
                continue
            margin = scores[j] - correct_class_score + 1 # note delta = 1
            if margin > 0:
                loss += margin
                W_der1 = X[i]
                dW[:, j] += W_der1 # j번째 열에 모두 더해주기
                dW[:, y[i]] -= W_der1
        # dW = dL/dW
```

```
# Right now the loss is a sum over all training examples, but we want it
# to be an average instead so we divide by num_train.
loss /= num_train
dW /= num_train
# Add regularization to the loss.
loss += reg * np.sum(W * W) # reg: regularization strength

#####
# T000:
# Compute the gradient of the loss function and store it dW.
# Rather than first computing the loss and then computing the derivative,
# it may be simpler to compute the derivative at the same time that the
# loss is being computed. As a result you may need to modify some of the
# code above to compute the gradient.
#####
# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
# loss function의 기울기 구하고 dW를 저장, loss는 숫자
pass
dW += 2 * reg * W # 뒷부분 더하기
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
```

03_svm_loss_vectorized(W, x, y, reg)



Generated by Notion to PDF

- 이 함수는 1번 함수에서 vectorized를 이용하였다.
- 똑같이 계산해주는데 `X.dot(w)`로 한꺼번에 계산해주고 `correct_score`도 `scores[range(num_train), y]`을 이용해서 한번에 슬라이싱 해준다.
- 그리고 maximum을 이용해서 배열에서 최대값을 한꺼번에 연산해준다.
- `correct_score`에 해당하는 부분도 한꺼번에 0으로 만들어준다.
- `dW`에 `correct_score`에 해당하는 부분을 0으로 지정해주었다.
- 그리고 `dW`는 `xi`들의 합으로 이루어지기 때문에 이를 이용하기 위해 새로운 `temp`이라는 배열을 선언해주었다.
- `temp[loss > 0] = 1`로 0보다 큰 부분에는 1을 할당해주었다.
- 그 후 `loss`가 0보다 큰 횟수만큼 `correct score`를 빼주어야 하므로 1의 개수를 새어주기 위해서 `np.sum(temp, axis = 1)`을 사용하였다./co

```

76 def svm_loss_vectorized(W, X, y, reg):
77     """
78     Structured SVM loss function, vectorized implementation.
79
80     Inputs and outputs are the same as svm_loss_naive.
81     """
82     loss = 0.0
83     dW = np.zeros(W.shape) # initialize the gradient as zero
84
85     #####
86     # TODO:
87     # Implement a vectorized version of the structured SVM loss, storing the
88     # result in loss.
89     #####
90     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
91     num_train = X.shape[0]
92     scores = X.dot(W)
93     # 슬라이싱
94     correct_class_score = scores[range(num_train), y].reshape(-1, 1)
95     #print(correct_class_score.shape) # 500 x 1
96     loss = np.maximum(0, scores - correct_class_score + 1) # 배열에서 최대값
97     loss[range(num_train), y] = 0
98     # dW: 3073 x 10
99     temp = np.zeros(loss.shape) # 500 x 10
100    temp[loss > 0] = 1
101    temp[range(num_train), y] -= np.sum(temp, axis = 1)
102    dW = np.dot(X.T, temp) # 3073 x 10, X에 있는 값을 모두 더해주면됨
103
104
105
106    dW = dW / num_train
107    loss = np.sum(loss)/num_train
108    loss += reg * np.sum(W * W)
109    pass
110

```

04_train



```

56     #####  

57     # TODO:  

58     # Sample batch_size elements from the training data and their  

59     # corresponding labels to use in this round of gradient descent.  

60     # Store the data in X_batch and their corresponding labels in  

61     # y_batch; after sampling X_batch should have shape (batch_size, dim)  

62     # and y_batch should have shape (batch_size, )  

63     #  

64     # Hint: Use np.random.choice to generate indices. Sampling with  

65     # replacement is faster than sampling without replacement.  

66     #####  

67     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  

68     x_idx = np.random.choice(num_train, batch_size, False)  

69     X_batch = X[x_idx]  

70     y_batch = y[x_idx] # index일치  

71     #print(X_batch.shape) # 200 x 3073  

72     pass  

73  

74     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  

75  

76     # evaluate loss and gradient  

77     loss, grad = self.loss(X_batch, y_batch, reg)  

78     loss_history.append(loss)  

79  

80     # perform parameter update  

81     #####  

82     # TODO:  

83     # Update the weights using the gradient and the learning rate.  

84     #####  

85     # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  

86     self.W -= learning_rate * grad  

87     pass  

88  

89     # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  

90  

91     if verbose and it % 100 == 0:  

92         print('iteration %d / %d: loss %f' % (it, num_iters, loss))  

93  

94     return loss_history
95

```

- 간단하게 W를 update해주는 부분만 추가되었다.

```

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  

for lr in learning_rates:  

    for rs in regularization_strengths:  

        svm = LinearSVM()  

        result = (lr, rs)  

        loss_hist = svm.train(X_train, y_train, learning_rate=lr, reg=rs,  

                              num_iters=1500, verbose=False)  

        y_train_pred = svm.predict(X_train)  

        training_accuracy = np.mean(y_train == y_train_pred)  

        y_val_pred = svm.predict(X_val)  

        validation_accuracy = np.mean(y_val == y_val_pred)  

  

        value = (training_accuracy, validation_accuracy)  

        results[result] = value  

        # 최적인지 판단  

        if (best_val < validation_accuracy):  

            best_val = validation_accuracy  

            best_svm = svm  

  

        pass  

  

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****  

  

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %f reg %f train accuracy: %f' % (lr, reg, train_accuracy))
    print('val accuracy: %f' % (val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)

```

▷ /content/drive/My Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_svm.py:110: RuntimeWarning: overflow encountered in double_scalars
loss += reg * np.sum(W ** 2)
/usr/local/lib/python3.8/dist-packages/numpy/core/fromnumeric.py:86: RuntimeWarning: overflow encountered in reduce
return ufunc.reduce(obj, axis, dtype, out, **passkwargs)
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_svm.py:110: RuntimeWarning: overflow encountered in multiply
loss += reg * np.sum(W ** 2)
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_svm.py:125: RuntimeWarning: overflow encountered in multiply
W += -2 * reg * W # 엇부분 대하기
/content/drive/My Drive/cs231n/assignments/assignment1/cs231n/classifiers/linear_classifier.py:86: RuntimeWarning: invalid value encountered in subtract
self.W -= learning_rate * grad
lr 1.000000e-07 reg 2.500000e-04 train accuracy: 0.369939 val accuracy: 0.386000
lr 1.000000e-07 reg 5.000000e-04 train accuracy: 0.354327 val accuracy: 0.361000
lr 5.000000e-05 reg 2.500000e-04 train accuracy: 0.054531 val accuracy: 0.050000
lr 5.000000e-05 reg 5.000000e+04 train accuracy: 0.100265 val accuracy: 0.087000
best validation accuracy achieved during cross-validation: 0.386000

- 모델을 만들어서 최적화를 시켜줬다. 각각의 parameter를 통해서 언제 최적의 값을 갖는지를 구했다.

1. Inline Question 2

Inline question 2

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way they do.

SVM weight의 이미지를 보면 원래의 이미지와 비슷한 것을 확인할 수 있다. 이는 내적했을 때 값이 크게 나와야 하므로 평행한 값으로 할당되었기 때문이다.

